AI ON THE RISE · AGENTS IN PRODUCTION

Without isolation, an agent in production becomes an operational risk

Microsoft launched containers with policies to contain agents

Victor Slompo · vssolutions.io

Swipe →

WHAT HAPPENED

Microsoft made agent containment part of the foundation of the stack

Microsoft Execution Containers became generally available as a policy-based execution layer. The idea is to limit files, network, interface, and process, with the policy outside the agent's control.

WHAT CHANGES

Production now requires native containment

01 The policy defines what the agent can access and do

02 Containment applies from the local device to the cloud

03 There are modes to observe, learn, and then enforce blocking

THE REAL PROBLEM IN COMPANIES

What I see in companies when an agent moves beyond the pilot

01 The agent gets too much access because no one modeled permissions properly

02 When something fails, there is no clear trail of what it tried to do

03 An improvised sandbox works in testing and breaks in real operations

I'VE ALREADY SOLVED THIS

I've already solved this in an enterprise account

01 Multi-agent platform in production in an enterprise account

02 US$ 5.5 million account

03 **About 60% efficiency gain and audited defect log**

MY TAKE

Measure instead of assume.

If agents are going to operate real systems and data, trust cannot be a promise. It has to start with containment, permission, and auditability. In practice, production requires evidence of what the agent did, tried to do, and where it was blocked.

WANT TO SEE HOW I DID IT?

Comment AGENTES.

I'll send you the case step by step by DM. I take AI off the drawing board and put it into production, with measured results.

Comment AGENTES · vssolutions.io