AI ON THE RISE · AGENT HARNESSES

The agent finished the task. The harness still failed.

A new benchmark separates task success from reliability.

Victor Slompo · vssolutions.io

Swipe →

WHAT HAPPENED

A paper scores the harness, not just the model

TraceTrial-Bench evaluates agent performance and harness reliability separately. It covers seven components: context, execution, governance, lifecycle, observability, tools and verification. It ran 4 harnesses, 9 models and 3,420 runs.

THE NUMBER

31.1%

of runs succeeded while failing at least one harness criterion

Source: Trial by Trace, 3,420 runs across 4 harnesses and 9 models.

THE REAL PROBLEM IN COMPANIES

What I see in companies: nobody can say where it failed

01 The demo passes, so the team assumes the agent is reliable.

02 Tool calls, retries and checks fail silently behind a correct final answer.

03 Without traces, every incident turns into a debate about the model.

I'VE ALREADY SOLVED THIS

My multi-agent platform at an enterprise account

01 A multi-agent platform running in production at a US$5.5 million account.

02 About 60% efficiency gain in production.

03 **Multi-agent in production requires an audited defect log.**

MY TAKE

Measure instead of assuming.

A correct final answer does not prove the agent worked. In practice, you need traces that show where each failure happened before you can fix it or trust it. The harness is where that evidence lives, so that is where I look first.

WANT TO SEE HOW I DID IT?

Comment AGENTS.

I'll send you the step-by-step of this case by DM. I take AI off the drawing board and put it into production, with measured results.

Comment AGENTS · vssolutions.io